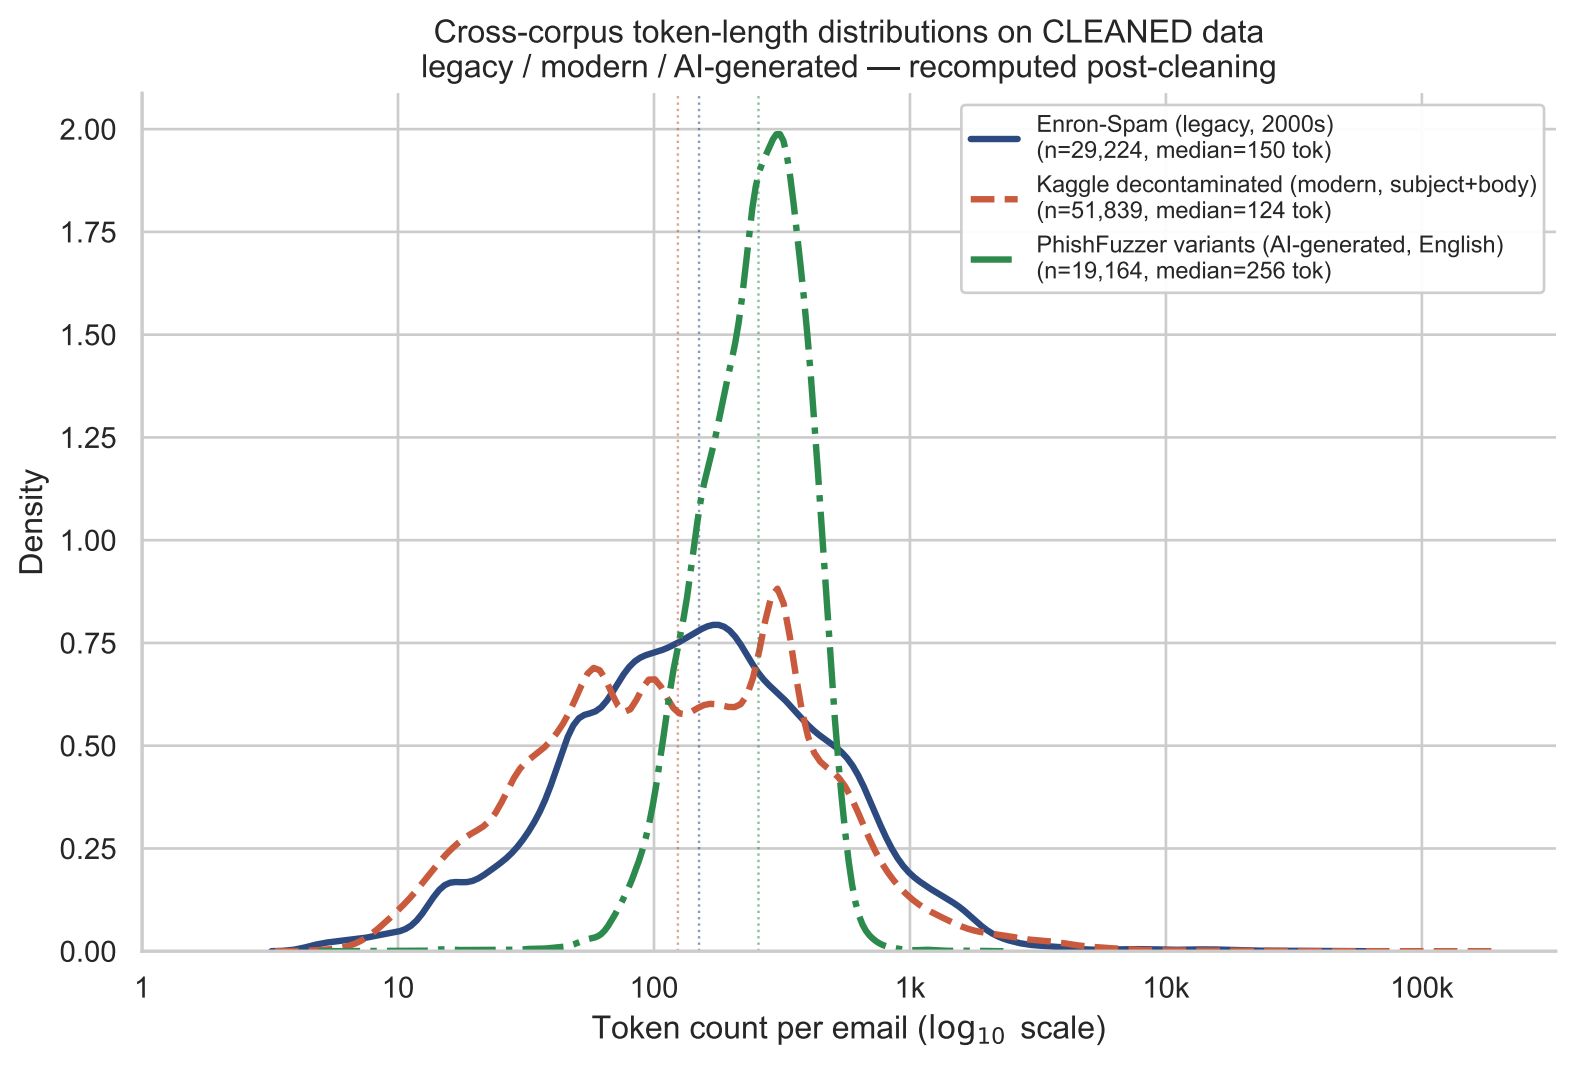
Cross-corpus token-length distributions on CLEANED data
legacy / modern / AI-generated — recomputed post-cleaning
Enron-Spam (legacy, 2000s)
(n=29,224, median=150 tok)
Kaggle decontaminated (modern, subject+body)
(n=51,839, median=124 tok)
PhishFuzzer variants (AI-generated, English)
(n=19,164, median=256 tok)
Density
2.00
1.75
1.50
1.25
1.00
0.75
0.50
0.25
0.00
1
10
100
1k
10k
100k
Token count per email (log10 scale)