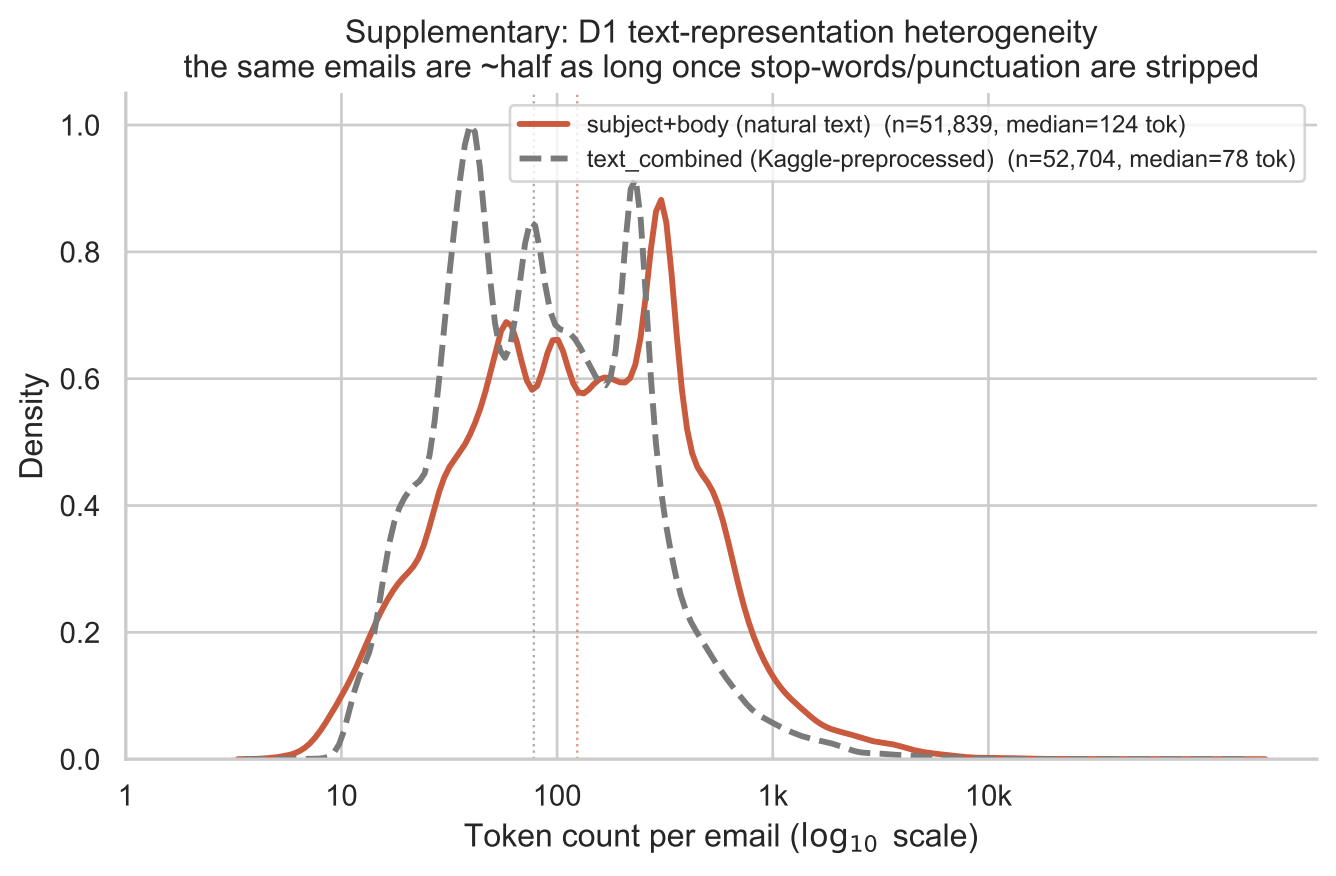
Supplementary: D1 text-representation heterogeneity
the same emails are ~half as long once stop-words/punctuation are stripped
subject+body (natural text) (n=51,839, median=124 tok)
text_combined (Kaggle-preprocessed) (n=52,704, median=78 tok)
Density
1.0
0.8
0.6
0.4
0.2
0.0
1
10
100
1k
10k
Token count per email (log10 scale)