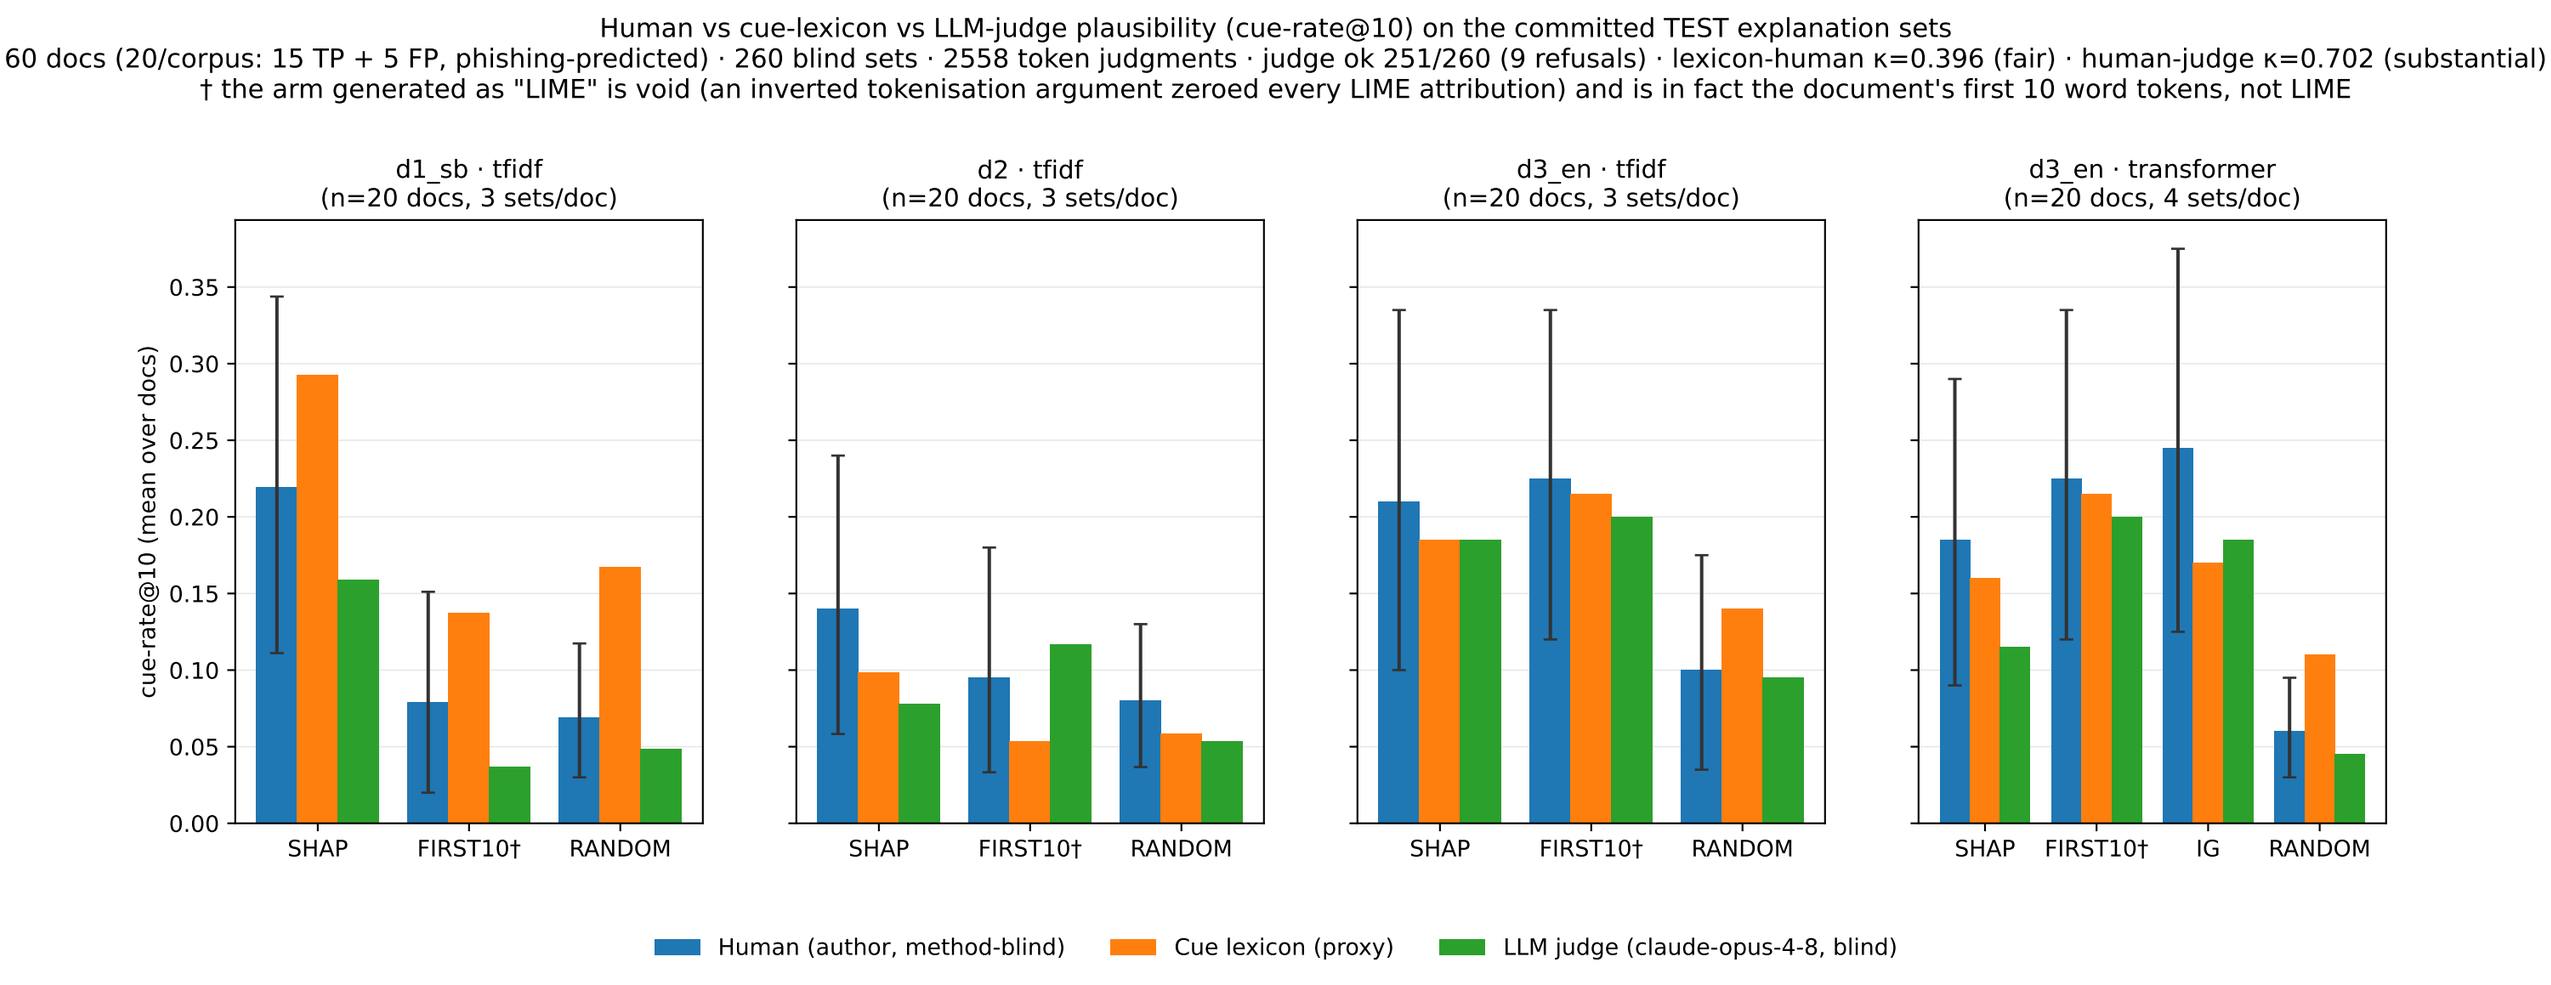
Human vs cue-lexicon vs LLM-judge plausibility (cue-rate@10) on the committed TEST explanation sets
60 docs (20/corpus: 15 TP + 5 FP, phishing-predicted) · 260 blind sets · 2558 token judgments · judge ok 251/260 (9 refusals) · lexicon-human κ=0.396 (fair) · human-judge κ=0.702 (substantial)
† the arm generated as "LIME" is void (an inverted tokenisation argument zeroed every LIME attribution) and is in fact the document's first 10 word tokens, not LIME
d1_sb · tfidf
(n=20 docs, 3 sets/doc)
d2 · tfidf
(n=20 docs, 3 sets/doc)
d3_en · tfidf
(n=20 docs, 3 sets/doc)
d3_en · transformer
(n=20 docs, 4 sets/doc)
cue-rate@10 (mean over docs)
0.00
0.05
0.10
0.15
0.20
0.25
0.30
0.35
SHAP
FIRST10†
RANDOM
IG
Human (author, method-blind)
Cue lexicon (proxy)
LLM judge (claude-opus-4-8, blind)